

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

Aim: To learn the Basics of Machine Learning

IDE: Google Colab

Theory:

Machine learning is a subset of artificial intelligence in the field of computer science that often uses statistical techniques to give computers the ability to "learn" (i.e., progressively improve performance on a specific task) with data, without being explicitly programmed. In the past decade, machine learning has given us self-driving cars, practical speech recognition, effective web search, and a vastly improved understanding of the human genome.

Machine learning tasks

Machine learning tasks are typically classified into two broad categories, depending on whether there is a learning "signal" or "feedback" available to a learning system:

- **Supervised learning:** The computer is presented with example inputs and their desired outputs, given by a "teacher", and the goal is to learn a general rule that maps inputs to outputs. As special cases, the input signal can be only partially available, or restricted to special feedback:
- **Semi-supervised learning:** the computer is given only an incomplete training signal: a training set with some (often many) of the target outputs missing.
- **Active learning:** the computer can only obtain training labels for a limited set of instances (based on a budget), and also has to optimize its choice of objects to acquire labels for. When used interactively, these can be presented to the user for labeling.
- **Reinforcement learning:** training data (in form of rewards and punishments) is given only as feedback to the program's actions in a dynamic environment, such as driving a vehicle or playing a game against an opponent.
- **Unsupervised learning:** No labels are given to the learning algorithm, leaving it on its own to find structure in its input. Unsupervised learning can be a goal in itself (discovering hidden patterns in data) or a means towards an end (feature learning).

Machine learning applications

In classification, inputs are divided into two or more classes, and the learner must produce a model that assigns unseen inputs to one or more (multi-label classification) of these classes. This is typically tackled in a supervised manner. Spam filtering is an example of classification, where the inputs are email (or other) messages and the classes are "spam" and "not spam". In regression, also a supervised problem, the outputs are continuous rather than discrete.

In clustering, a set of inputs is to be divided into groups. Unlike in classification, the groups are not known beforehand, making this typically an unsupervised task. Density estimation finds the distribution of inputs in

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

some space. Dimensionality reduction simplifies inputs by mapping them into a lower- dimensional space. Topic modeling is a related problem, where a program is given a list of human language documents and is tasked with finding out which documents cover similar topics.

Machine learning Approaches

Decision tree learning: Decision tree learning uses a decision tree as a predictive model, which maps observations about an item to conclusions about the item's target value. Association rule learning Association rule learning is a method for discovering interesting relations between variables in large databases.

Artificial neural networks

An artificial neural network (ANN) learning algorithm, usually called "neural network" (NN), is a learning algorithm that is vaguely inspired by biological neural networks. Computations are structured in terms of an interconnected group of artificial neurons, processing information using a connectionist approach to computation. Modern neural networks are non-linear statistical data modeling tools. They are usually used to model complex relationships between inputs and outputs, to find patterns in data, or to capture the statistical structure in an unknown joint probability distribution between observed variables.

Deep learning

Falling hardware prices and the development of GPUs for personal use in the last few years have contributed to the development of the concept of deep learning which consists of multiple hidden layers in an artificial neural network. This approach tries to model the way the human brain processes light and sound into vision and hearing. Some successful applications of deep learning are computer vision and speech recognition.

Inductive logic programming

Inductive logic programming (ILP) is an approach to rule learning using logic programming as a uniform representation for input examples, background knowledge, and hypotheses. Given an encoding of the known background knowledge and a set of examples represented as a logical database of facts, an ILP system will derive a hypothesized logic program that entails all positive and no negative examples. Inductive programming is a related field that considers any kind of programming languages for representing hypotheses (and not only logic programming), such as functional programs.

Support vector machines

Support vector machines (SVMs) are a set of related supervised learning methods used for classification and regression. Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that predicts whether a new example falls into one category or the other.

Clustering

Cluster analysis is the assignment of a set of observations into subsets (called clusters) so that observations within the same cluster are similar according to some pre designated criterion or criteria, while observations drawn from different clusters are dissimilar. Different clustering techniques make different assumptions on the

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

structure of the data, often defined by some similarity metric and evaluated for example by internal compactness (similarity between members of the same cluster) and separation between different clusters. Other methods are based on estimated density and graph connectivity. Clustering is a method of unsupervised learning, and a common technique for statistical data analysis.

Bayesian networks

A Bayesian network, belief network or directed acyclic graphical model is a probabilistic graphical model that represents a set of random variables and their conditional independencies via a directed acyclic graph (DAG). For example, a Bayesian network could represent the probabilistic relationships between diseases and symptoms. Given symptoms, the network can be used to compute the probabilities of the presence of various diseases. Efficient algorithms exist that perform inference and learning.

Reinforcement learning

Reinforcement learning is concerned with how an agent ought to take actions in an environment so as to maximize some notion of long-term reward. Reinforcement learning algorithms attempt to find a policy that maps states of the world to the actions the agent ought to take in those states. Reinforcement learning differs from the supervised learning problem in that correct input/output pairs are never presented, nor sub-optimal actions explicitly corrected.

Similarity and metric learning

In this problem, the learning machine is given pairs of examples that are considered similar and pairs of less similar objects. It then needs to learn a similarity function (or a distance metric function) that can predict if new objects are similar. It is sometimes used in Recommendation systems.

Genetic algorithms

A genetic algorithm (GA) is a search heuristic that mimics the process of natural selection, and uses methods such as mutation and crossover to generate new genotype in the hope of finding good solutions to a given problem. In machine learning, genetic algorithms found some uses in the 1980s and 1990s. Conversely, machine learning techniques have been used to improve the performance of genetic and evolutionary algorithms.

Rule-based machine learning

Rule-based machine learning is a general term for any machine learning method that identifies, learns, or evolves "rules" to store, manipulate or apply, knowledge. The defining characteristic of a rule-based machine learner is the identification and utilization of a set of relational rules that collectively represent the knowledge captured by the system. This is in contrast to other machine learners that commonly identify a singular model that can be universally applied to any instance in order to make a prediction. Rule-based machine learning approaches include learning classifier systems, association rule learning, and artificial immune systems.

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

Feature selection approach

Feature selection is the process of selecting an optimal subset of relevant features for use in model construction. It is assumed the data contains some features that are either redundant or irrelevant, and can thus be removed to reduce calculation cost without incurring much loss of information. Common optimality criteria include accuracy, similarity and information measures.

Pre Lab Exercise:

a. **Write 5 applications that you encounter daily, where Machine Learning is used?**

1. Recommendation Systems (Amazon, Netflix, Flipkart).
2. Social Media Connections & Feeds (Facebook, Youtube, LinkedIn).
3. Virtual Personal Assistants (Siri, Alexa, Google Assistant).
4. Fraud Detection in Banking & Payments (Credit card transactions, UPI apps).
5. Image & Facial Recognition (Smartphone Face ID, Google Photos, law enforcement tools).

b. **What do you mean by Supervised Learning?**

Supervised learning is a technique where the algorithm is trained using labeled data. Each training example consists of an input and the corresponding correct output. The model learns the relationship between inputs and outputs, and once trained, it can predict outcomes for new, unseen inputs.

- **Key idea:** Learn from examples with known answers.
- **Applications:** Spam email detection, medical diagnosis, stock price prediction.

c. **What do you mean by Unsupervised Learning?**

Unsupervised learning deals with unlabeled data, meaning the system is not told the correct output. Instead, the algorithm tries to discover hidden patterns, structures, or groupings within the data. It is often used for clustering and dimensionality reduction.

- **Key idea:** Find structure in data without guidance.
- **Applications:** Customer segmentation, grouping similar documents, anomaly detection.

d. **What do you mean by Reinforcement Learning?**

Reinforcement learning is based on the concept of an agent interacting with an environment. The agent performs actions, receives feedback in the form of rewards or penalties, and learns strategies to maximize cumulative rewards over time. It balances exploration (trying new actions) and exploitation (using known strategies).

- **Key idea:** Learn by trial and error with rewards.
- **Applications:** Game-playing AI, robotics, self-driving cars, adaptive cybersecurity systems.

Program (Code):



Perform the following tasks:

1. Load a dataset in your IDE
2. Observe the statistics of all the features
3. Obtain the shape of the dataset
4. Separate all the features


 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

5. Fill the missing values, if any, using the statistically relevant value
6. Observe the Box-Plot of each feature
7. Comment on the IQR range and outliers



Results:




 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059




 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059



 Marwadi University	Marwadi University	
	Faculty of Technology	
	Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059



 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059



Observation and Result Analysis:

A. Statistics of the dataset

From the statistical summary obtained using the `describe()` function, we observed the count, mean, standard deviation, minimum value, maximum value, and quartile values (25%, 50%, and 75%) for all numerical features. The count value shows that most of the features have an equal number of records, which means there are very few or no missing values. The mean and median (50%) values help in understanding the central tendency of the data. If the mean and median values are close, the data is normally distributed. The standard deviation indicates how much the data is spread from the mean. Features with higher standard deviation show more variation in data. The minimum and maximum values give the range of the data. Overall, the statistical summary helps to understand the distribution, spread, and nature of the dataset before further processing.

B. IQR Outliers

From the boxplot of each numerical feature, the Inter Quartile Range (IQR) was observed. The IQR is calculated as the difference between the third quartile (Q3) and the first quartile (Q1), i.e., $IQR = Q3 - Q1$. It represents the middle 50% of the data.

Using the IQR method, values below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$ are considered outliers. In the boxplot, these outliers are shown as points outside the whiskers.

From the observed boxplots, some features showed very few outliers, while some features had no significant outliers. The IQR range helped in understanding the spread and variability of the data. Features with larger IQR indicate more spread in the middle portion of the dataset.

 Marwadi University	Marwadi University Faculty of Technology Department of Information and Communication Technology	
Subject: Artificial Intelligence (01CT0616)	Aim: To learn the Basics of Machine Learning	
Experiment No: 01	Date: 10/2/2026	Enrollment No: 92301733059

Post Lab Exercise:

a. Which is a better mean of replacing the missing value-mean, median or mode?

The better method of replacing missing values depends on the type and distribution of data. For numerical data, mean or median can be used. If the data does not contain outliers and is normally distributed, mean is suitable. If the data contains outliers, the median is better because it is not affected much by extreme values.

For categorical data, mode is generally used because it replaces the missing value with the most frequently occurring category.

b. How does IQR range help to understand the nature of the distribution of the data?

The IQR (Inter Quartile Range) represents the difference between Q3 and Q1 and shows the spread of the middle 50% of the data. It helps to understand how the data is distributed and whether the data is spread out or closely packed. IQR is also useful in detecting outliers. Values lying below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$ are considered outliers. Therefore, IQR helps in understanding the variability and the presence of extreme values in the dataset.

c. Which are the ways to replace the categorical data to the numerical ones?

Categorical data can be converted into numerical form using different encoding techniques. One common method is assigning numerical labels to each category. Another method is creating separate binary columns for each category (dummy variables). These methods help in converting categorical data into numerical form so that machine learning algorithms can process the data properly.

Conclusion:

In this experiment, the basic concepts of Machine Learning were studied and implemented using Google Collab. The dataset was successfully loaded and analyzed. By observing its statistical summary and shape. Missing values were identified and handled using appropriate statistical methods. Box plots were also generated to visualize the distribution of features and to detect possible outliers using the IQR method. Overall, this experiment helped in understanding the initial steps involved in data analysis and preprocessing before applying machine learning models.